


CS5182 Computer Graphics Introduction to Tensorflow & 3D Object Detection Tutorial




2024/25 Semester A

City University of Hong Kong (DG)

Deep Learning Framework

- Why we need framework for coding deep learning?
 - Easy implementation of neural network, e.g. layers, loss functions ...
 - Automatically calculate gradients to update models
 - Fast integration, especially using GPU
 - Community are using them, lot of resources, e.g. GitHub
- TensorFlow 1.0 v.s. PyTorch


	Developed by	Graphs	Difficulty	Speed
 TensorFlow	Google	Static graphs	Hard	Fast
PYTORCH	Facebook	Dynamic graphs	Easy	Fast enough

- Keras, TFLearn: high-level API built upon TensorFlow
- TensorFlow 2.0: using dynamic graphs, but totally different (less resource)

Tensor + Flow

□ Main components of TensorFlow

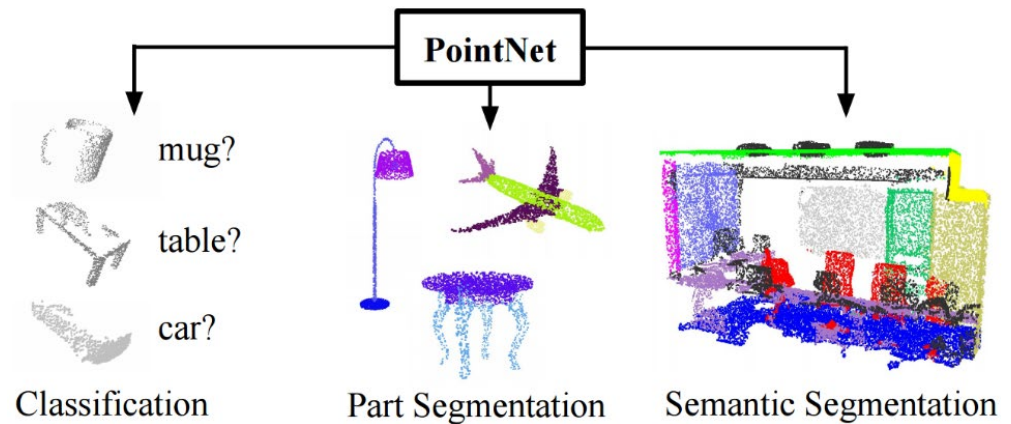
- Tensor: Define variables, including input/output, model weights.
- Graph: Define calculation. Build static graph to connect those tensors.
- Session: Execute calculation with data feeding. (Tensors flow in graph!)



Tensor + Flow

□ Two examples

- MNIST digit recognition using 2D CNN
- 3D point cloud classification using PointNet



First step: Install TensorFlow

- ❑ Preliminary: Python, basic Linux, GPU machine
- ❑ Step1: Create a virtual environment (recommend Anaconda)

```
(base) qianyue@ss4028b:~$ conda create --name tf_env
...
(base) qianyue@ss4028b:~$ conda activate tf_env
(tf_env) qianyue@ss4028b:~$ █
```

- ❑ Step2: Install TensorFlow in new environment (recommend tf1)


```
(tf_env) qianyue@ss4028b:~$ conda install tensorflow-gpu
(tf_env) qianyue@ss4028b:~$ conda install tensorflow-gpu==1.9.0
```

- ❑ Step3: Check whether success installed

```
>>> import tensorflow as tf
>>> tf.__version__
'1.9.0'
>>> █
```

MNIST digit recognition using 2D CNN

- Example 1: MNIST digit recognition using 2D CNN
 - Input: 28x28 BW images
 - Output: which 0~9 digit?
 - Dataset: MNIST, 60k training, 10k testing with ground-truth label provided.
- Develop a simple CNN model for MNIST
 - Load dataset
 - Build neural network
 - Define loss function and learning strategies
 - Training iteratively
 - Evaluate performance and save model
 - Use TensorBoard to monitor training process




1. Load dataset

□ 1.1 Load MNIST

```
11 import tensorflow as tf
12
13 # Import MNIST data
14 from tensorflow.examples.tutorials.mnist import input_data
15 mnist = input_data.read_data_sets("/tmp/data/", one_hot=True)
16
17 # batch_x, batch_y = mnist.train.next_batch(batch_size)
18 # test_x, test_y = mnist.test.images[:20], mnist.test.labels[:20]
```

□ 1.2 Define placeholder for input and ground truth label


```
24 num_input = 784 # MNIST data input (img shape: 28*28)
25 num_classes = 10 # MNIST total classes (0-9 digits)
26
27 # tf Graph input
28 X = tf.placeholder(tf.float32, [None, num_input])
29 Y = tf.placeholder(tf.float32, [None, num_classes])
30
31 keep_prob = tf.placeholder(tf.float32) # dropout (keep probability)
32 is_reuse = tf.placeholder(tf.bool, shape=())
```



2. Build neural network

2.1 Define network

```
39 # Create the neural network
40 def conv_net(x, dropout, reuse):
41     n_classes = 10
42     # Define a scope for reusing the variables
43     with tf.variable_scope('ConvNet', reuse=reuse):
44         # MNIST data input is a 1-D vector of 784 features (28*28 pixels)
45         # Reshape to match picture format [Height x Width x Channel]
46         # Tensor input become 4-D: [Batch Size, Height, Width, Channel]
47         x = tf.reshape(x, shape=[-1, 28, 28, 1])
48         # Convolution Layer with 32 filters and a kernel size of 5
49         conv1 = tf.layers.conv2d(x, 32, 5, activation=tf.nn.relu)
50         # Max Pooling (down-sampling) with strides of 2 and kernel size of 2
51         conv1 = tf.layers.max_pooling2d(conv1, 2, 2)
52         # Convolution Layer with 64 filters and a kernel size of 3
53         conv2 = tf.layers.conv2d(conv1, 64, 3, activation=tf.nn.relu)
54         # Max Pooling (down-sampling) with strides of 2 and kernel size of 2
55         conv2 = tf.layers.max_pooling2d(conv2, 2, 2)
56
57         # Flatten the data to a 1-D vector for the fully connected layer
58         fc1 = tf.contrib.layers.flatten(conv2)
59         # Fully connected layer (in tf contrib folder for now)
60         fc1 = tf.layers.dense(fc1, 1024)
61         # Apply Dropout (if is_training is False, dropout is not applied)
62         fc1 = tf.nn.dropout(fc1, dropout)
63
64         # Output layer, class prediction
65         out = tf.layers.dense(fc1, n_classes)
66     return out
```



Build neural network & Loss function


□ 2.2 From placeholder X to prediction

```
33 # tf Graph input
34 X = tf.placeholder(tf.float32, [None, num_input])
35 Y = tf.placeholder(tf.float32, [None, num_classes])
36 keep_prob = tf.placeholder(tf.float32) # dropout (keep probability)
37 is_reuse = tf.placeholder(tf.bool, shape=())

131 logits = conv_net(X, keep_prob, is_reuse)
132 prediction = tf.nn.softmax(logits)
```

□ 3. Define loss function

```
100 # Define loss and optimizer
101 loss_op = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(
102     logits=logits, labels=Y))
```



4. Model Training


□ 4.1 Define training hyper parameters

```
17 # Training Parameters
18 learning_rate = 0.001
19 num_steps = 200
20 batch_size = 128

103 optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate)
104 train_op = optimizer.minimize(loss_op)
```

□ 4.2 Train iteratively

```
145 # Initialize the variables (i.e. assign their default value)
146 init = tf.global_variables_initializer()
147 # Start training
148 with tf.Session() as sess:
149     # Run the initializer
150     sess.run(init)
151     for step in range(1, num_steps+1):
152         batch_x, batch_y = mnist.train.next_batch(batch_size)
153         # Run optimization op (backprop)
154         sess.run(train_op, feed_dict={X: batch_x, Y: batch_y, keep_prob: 0.8, is_reuse: True})
155         # Calculate batch loss and accuracy
156         loss, acc = sess.run([loss_op, accuracy], feed_dict={X: batch_x,
157                                                             Y: batch_y,
158                                                             keep_prob: 1.0,
159                                                             is_reuse: False})
160         print("Step " + str(step) + ", Minibatch Loss= " + \
161               "{:.4f}".format(loss) + ", Training Accuracy= " + \
162               "{:.3f}".format(acc))
163
164     print("Optimization Finished!")
```



5. Evaluate performance and save model

5.1 Save trained model in path


```
192 saver = tf.train.Saver()
193 with tf.Session() as sess:
194     saver.save(sess, path, global_step=step)
```

5.2 Define accuracy

```
107 # Evaluate model
108 correct_pred = tf.equal(tf.argmax(prediction, 1), tf.argmax(Y, 1))
109 accuracy = tf.reduce_mean(tf.cast(correct_pred, tf.float32))
```

5.3 Evaluate on testing dataset

```
166 with tf.Session() as sess:
167     # Calculate accuracy for 256 MNIST test images
168     print("Testing Accuracy:", \
169         sess.run(accuracy, feed_dict={X: mnist.test.images[:256],
170                                         Y: mnist.test.labels[:256],
171                                         keep_prob: 1.0,
172                                         is_reuse: False}))
---
```



6. Use TensorBoard to monitor training process

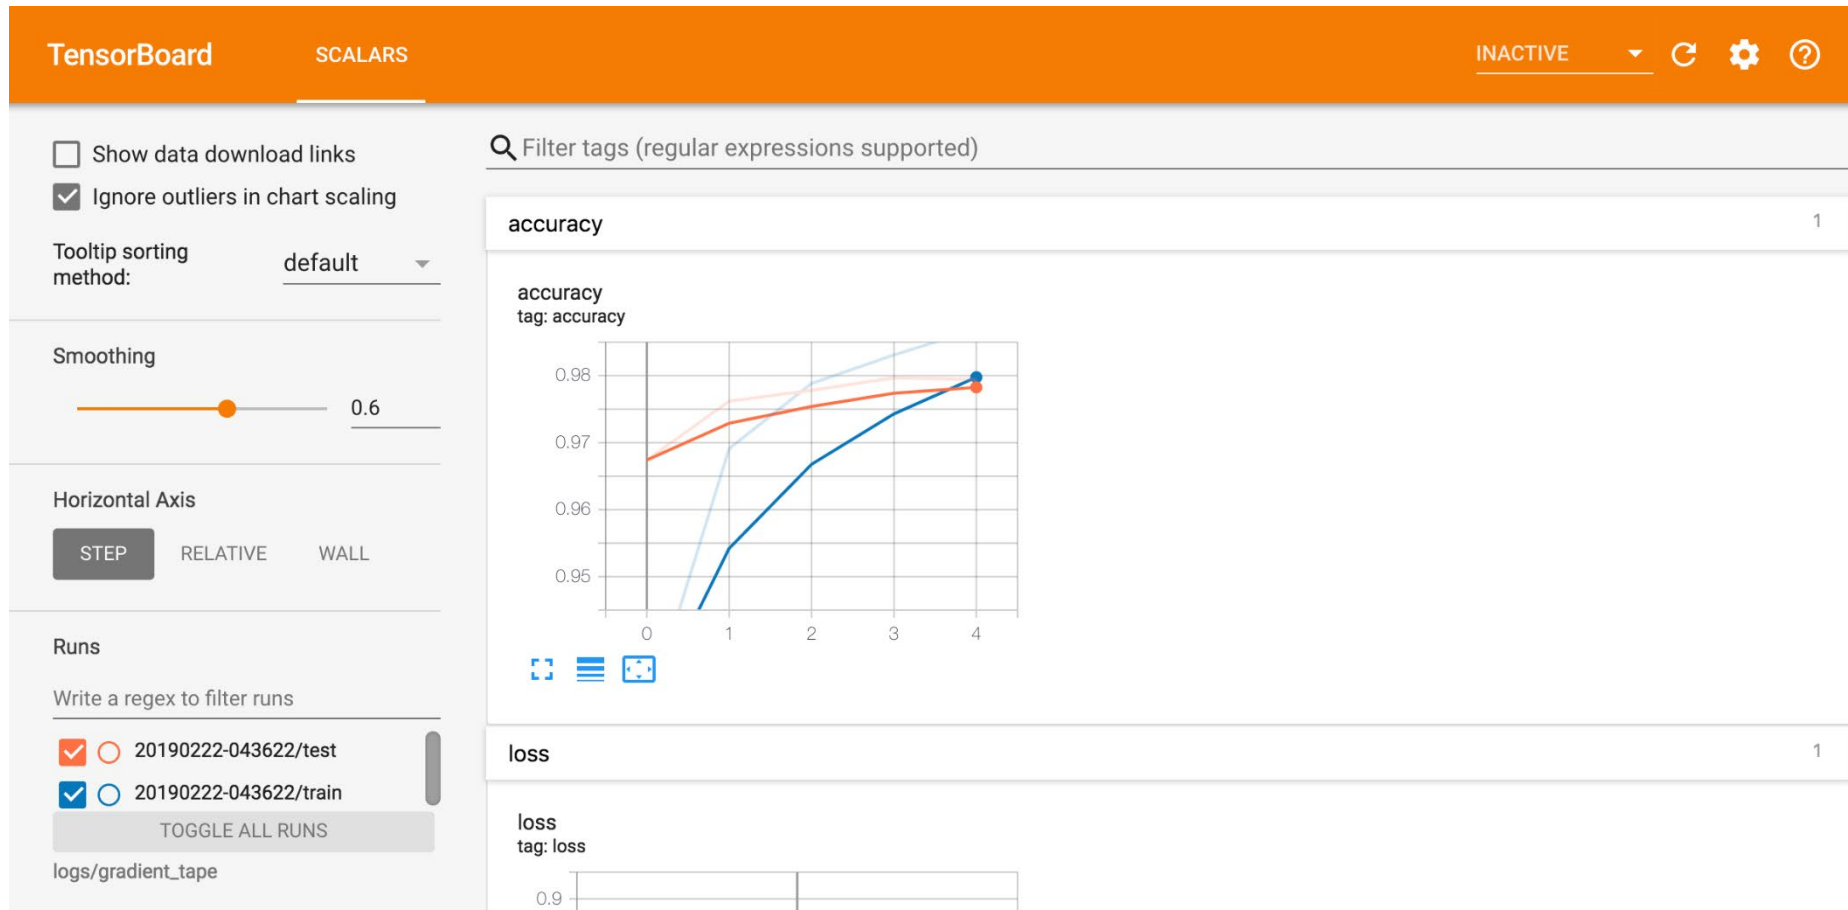
□ 3.1 Record information

```
54 # Create a summary to monitor cost tensor
55 tf.summary.scalar("loss", cost)
56 # Create a summary to monitor accuracy tensor
57 tf.summary.scalar("accuracy", acc)
58 # Merge all summaries into a single op
59 merged_summary_op = tf.summary.merge_all()
60
61 logs_path = '/tmp/tensorflow_logs/example/'
62 # Start training
63 with tf.Session() as sess:
64     sess.run(init)
65     # op to write logs to Tensorboard
66     summary_writer = tf.summary.FileWriter(logs_path, graph=tf.get_default_graph())
67     for step in range(1, num_steps+1):
68         batch_x, batch_y = mnist.train.next_batch(batch_size)
69         # Run optimization op (backprop), cost op (to get loss value)
70         # and summary nodes
71         _, summary = sess.run([train_op, merged_summary_op],
72                               feed_dict={x: batch_x, y: batch_y, keep_prob: 0.8})
73         # Write logs at every iteration
74         summary_writer.add_summary(summary, step)
75     --
```

□ 3.2 Visualize recorded results


```
(tf_env) qian Yue@ss4028b:~$ tensorboard --logdir='/tmp/tensorflow_logs/example/'
```

TensorBoard demo



3D point cloud classification using PointNet

□ Point cloud v.s. image



V.S.



**Point Cloud
Representation**

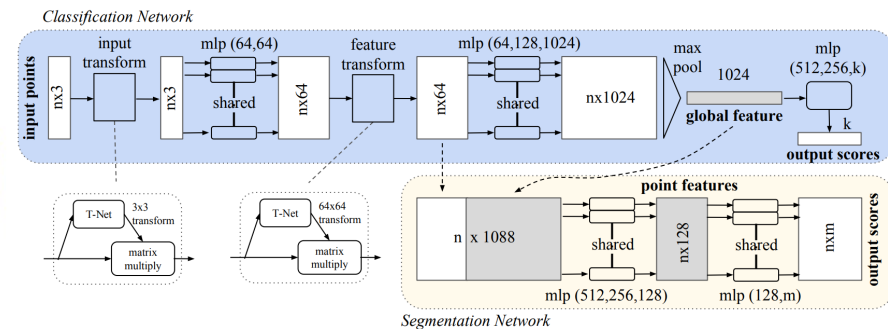
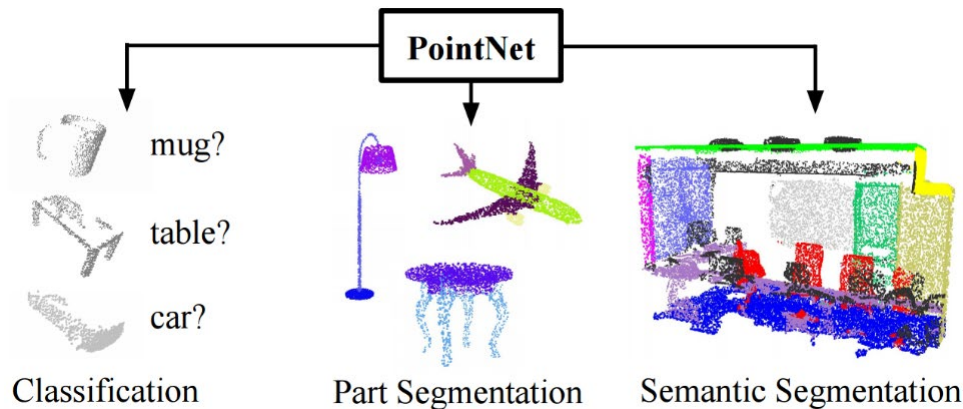


Bunny 3D Model

- Size: $[H, W, 3]$, where feature is (r, g, b)
- 2D regular structure, encoded as matrix
- Size: $[N, 3]$
- Feature is (x, y, z)
- 3D irregular structure, discrete representation for 3D shape

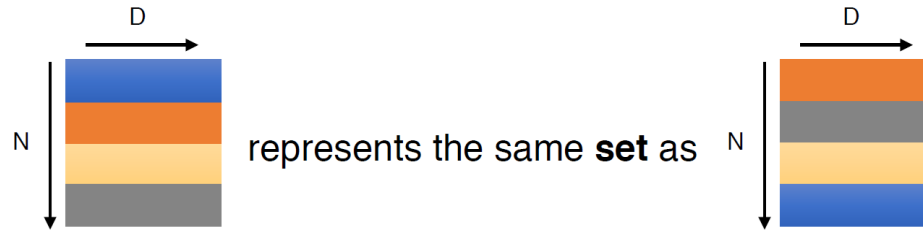
3D point cloud classification using PointNet

- PointNet: Deep Learning on Point Sets for 3D Classification and Segmentation ([paper](#))
- Input: $[N, 3]$ point clouds
- Output:
 - Classification: $[K, 1]$ as scores for K classification classes
 - Segmentation: $[N, M]$ as point-wise scores for M segmentation classes



3D point cloud classification using PointNet

- Challenge1: Permutation Invariance: Point cloud is a set of unordered points




Model needs to be invariant to $N!$ permutations

- Solution: Symmetric Function

- Examples:

$$f(x_1, x_2, \dots, x_n) = \max\{x_1, x_2, \dots, x_n\}$$


$$f(x_1, x_2, \dots, x_n) = x_1 + x_2 + \dots + x_n$$



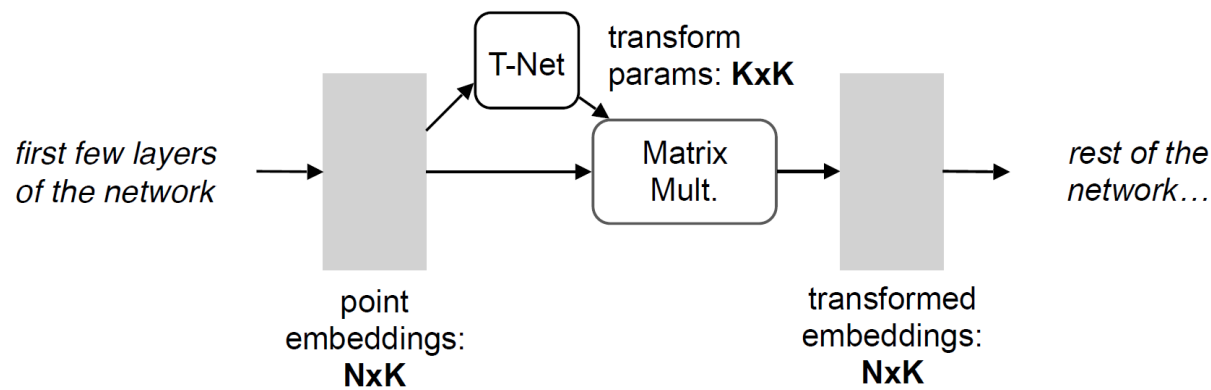
Construction Symmetric Functions by Neural Networks

3D point cloud classification using PointNet

- Challenge2: Invariance under geometric transformations: Point cloud rotation should not affect classification results.



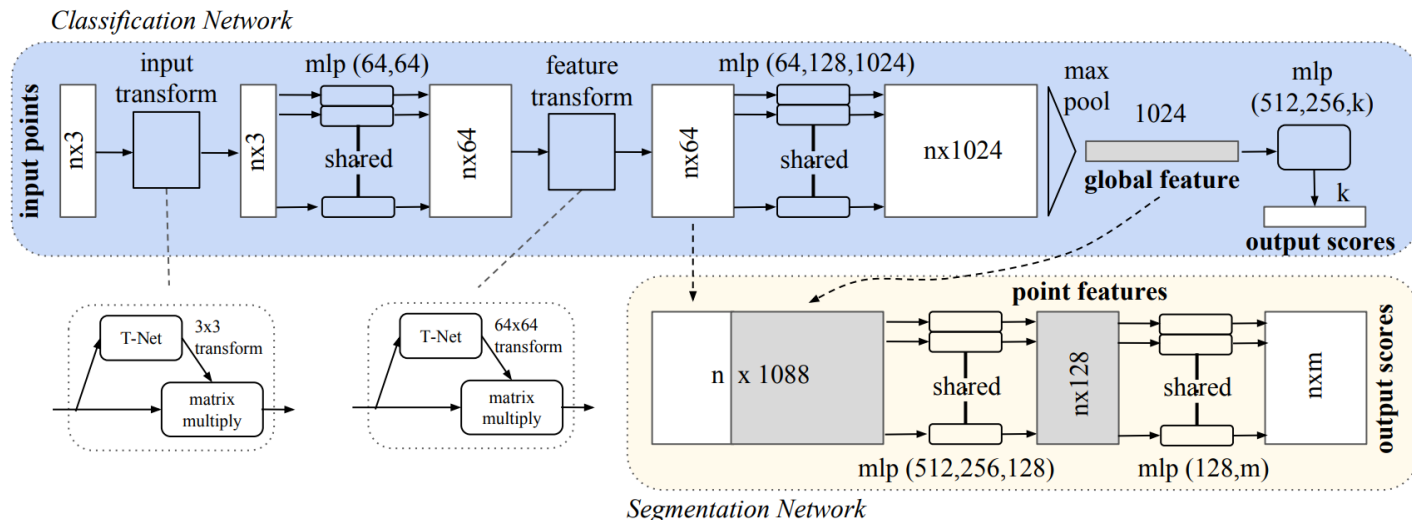
- Solution: Data dependent transformation for automatic alignment
 - Matrix multiplication



Embedding Space Alignment

3D point cloud classification using PointNet

- Input is a set of point cloud data, represented as an $n \times 3$ 2D tensor, where n represents the number of point clouds, and 3 corresponds to the xyz coordinates.
- The input data is first aligned by multiplying it with a transformation matrix learned by **T-Net**, which ensures the invariance of the model to specific spatial transformations.
- After extracting the features of each point cloud data through multiple MLPs, use a **T-Net** to align the features.
- Perform a **max-pooling** operation on each dimension of the feature to obtain the final **global feature**.
- Output
 - For the classification task, the global feature is used to predict the final classification score through MLPs;
 - For the segmentation task, the global feature and the local features of each point cloud learned before are concatenated, and then the classification result of each data point is obtained through MLPs .



PointNet code (main - 1/3)

```
1 import argparse
2 import math
3 import h5py
4 import numpy as np
5 import tensorflow as tf
6 import socket
7 import importlib
8 import os
9 import sys
10 BASE_DIR = os.path.dirname(os.path.abspath(__file__))
11 sys.path.append(BASE_DIR)
12 sys.path.append(os.path.join(BASE_DIR, 'models'))
13 sys.path.append(os.path.join(BASE_DIR, 'utils'))
14 import provider
15 import tf_util

16 parser = argparse.ArgumentParser()
17 parser.add_argument('--gpu', type=int, default=0, help='GPU to use [default: GPU 0]')
18 parser.add_argument('--model', default='pointnet_cls', help='Model name: pointnet_cls or pointnet_cls_basic [default: pointnet_cls]')
19 parser.add_argument('--log_dir', default='log', help='Log dir [default: log]')
20 parser.add_argument('--num_point', type=int, default=1024, help='Point Number [256/512/1024/2048] [default: 1024]')
21 parser.add_argument('--max_epoch', type=int, default=250, help='Epoch to run [default: 250]')
22 parser.add_argument('--batch_size', type=int, default=32, help='Batch Size during training [default: 32]')
23 parser.add_argument('--learning_rate', type=float, default=0.001, help='Initial learning rate [default: 0.001]')
24 parser.add_argument('--momentum', type=float, default=0.9, help='Initial learning rate [default: 0.9]')
25 parser.add_argument('--optimizer', default='adam', help='adam or momentum [default: adam]')
26 parser.add_argument('--decay_step', type=int, default=200000, help='Decay step for lr decay [default: 200000]')
27 parser.add_argument('--decay_rate', type=float, default=0.7, help='Decay rate for lr decay [default: 0.8]')
28 FLAGS = parser.parse_args()
29
30 BATCH_SIZE = FLAGS.batch_size
31 NUM_POINT = FLAGS.num_point
32 MAX_EPOCH = FLAGS.max_epoch
33 BASE_LEARNING_RATE = FLAGS.learning_rate
34 GPU_INDEX = FLAGS.gpu
35 MOMENTUM = FLAGS.momentum
36 OPTIMIZER = FLAGS.optimizer
37 DECAY_STEP = FLAGS.decay_step
38 DECAY_RATE = FLAGS.decay_rate
39
40 MODEL = importlib.import_module(FLAGS.model) # import network module
41 MODEL_FILE = os.path.join(BASE_DIR, 'models', FLAGS.model+'.py')
42 LOG_DIR = FLAGS.log_dir
43 ..
```

Install TensorFlow, h5py
Import all the dependencies

provider: Load data and perform data processing
tf_util: Define layers (convolution, FC, dropout) for point clouds

classification model

Training hyper-parameters

PointNet code (main - 2/3)

```
73 def get_learning_rate(batch):
74     learning_rate = tf.train.exponential_decay(
75         BASE_LEARNING_RATE, # Base learning rate.
76         batch * BATCH_SIZE, # Current index into the dataset.
77         DECAY_STEP,         # Decay step.
78         DECAY_RATE,         # Decay rate.
79         staircase=True)
80     learning_rate = tf.maximum(learning_rate, 0.00001) # CLIP THE LEARNING RATE!
81     return learning_rate
82
83 def get_bn_decay(batch):
84     bn_momentum = tf.train.exponential_decay(
85         BN_INIT_DECAY,
86         batch*BATCH_SIZE,
87         BN_DECAY_DECAY_STEP,
88         BN_DECAY_DECAY_RATE,
89         staircase=True)
90     bn_decay = tf.minimum(BN_DECAY_CLIP, 1 - bn_momentum)
91     return bn_decay
92
```

Training hyper-parameters

Run on specific GPU

```
93 def train():
94     with tf.Graph().as_default():
95         with tf.device('/gpu:' + str(GPU_INDEX)):
96             pointclouds_pl, labels_pl = MODEL.placeholder_inputs(BATCH_SIZE, NUM_POINT)
97             is_training_pl = tf.placeholder(tf.bool, shape=())
98             print(is_training_pl)
99
100             # Note the global_step=batch parameter to minimize.
101             # That tells the optimizer to helpfully increment the 'batch' parameter for you every time it trains.
102             batch = tf.Variable(0)
103             bn_decay = get_bn_decay(batch)
104             tf.summary.scalar('bn_decay', bn_decay)
105
106             # Get model and loss
107             pred, end_points = MODEL.get_model(pointclouds_pl, is_training_pl, bn_decay=bn_decay)
108             loss = MODEL.get_loss(pred, labels_pl, end_points)
109             tf.summary.scalar('loss', loss)
110
111             correct = tf.equal(tf.argmax(pred, 1), tf.to_int64(labels_pl))
112             accuracy = tf.reduce_sum(tf.cast(correct, tf.float32)) / float(BATCH_SIZE)
113             tf.summary.scalar('accuracy', accuracy)
114
115             # Add ops to save and restore all the variables.
116             saver = tf.train.Saver()
```

```
12 def placeholder_inputs(batch_size, num_point):
13     pointclouds_pl = tf.placeholder(tf.float32, shape=(batch_size, num_point, 3))
14     labels_pl = tf.placeholder(tf.int32, shape=(batch_size))
15     return pointclouds_pl, labels_pl
```

Call PointNet network
pred: [B, K]
classification scores

Define loss

Define
accuracy

In order to store trained
model

PointNet code (main - 3/3)

```
127 # Create a session
128 config = tf.ConfigProto()
129 config.gpu_options.allow_growth = True
130 config.allow_soft_placement = True
131 config.log_device_placement = False
132 sess = tf.Session(config=config)
133
134 # Add summary writers
135 #merged = tf.merge_all_summaries()
136 merged = tf.summary.merge_all()
137 train_writer = tf.summary.FileWriter(os.path.join(LOG_DIR, 'train'),
138                                     sess.graph)
139 test_writer = tf.summary.FileWriter(os.path.join(LOG_DIR, 'test'))
140
141 # Init variables
142 init = tf.global_variables_initializer()
143 sess.run(init, {is_training_pl: True})
144
145 for epoch in range(MAX_EPOCH):
146     log_string('*** EPOCH %03d ***' % (epoch))
147     sys.stdout.flush()
148
149     train_one_epoch(sess, ops, train_writer)
150     eval_one_epoch(sess, ops, test_writer)
151
152 # Save the variables to disk.
153 if epoch % 10 == 0:
154     save_path = saver.save(sess, os.path.join(LOG_DIR, "model.ckpt"))
155     log_string("Model saved in file: %s" % save_path)
```

Create session

For TensorBoard visualization

Initialize variables

Train iteratively
Epoch: train whole dataset once

Store model

Network code (pointnet_cls.get_model)

```
9 import tf_util
10 from transform_nets import input_transform_net, feature_transform_net
```

```
11 def get_model(point_cloud, is_training, bn_decay=None):
12     """ Classification PointNet, input is BxNx3, output Bx40 """
13     batch_size = point_cloud.get_shape()[0].value
14     num_point = point_cloud.get_shape()[1].value
15     end_points = {}
```

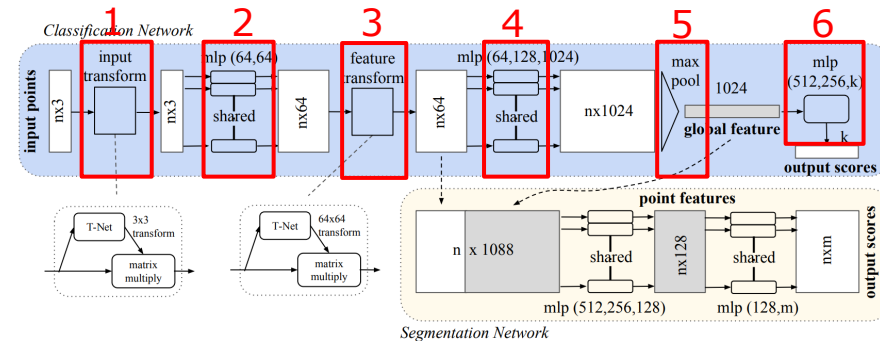
```
16     with tf.variable_scope('transform_net1') as sc:
17         transform = input_transform_net(point_cloud, is_training, bn_decay, K=3)
18         point_cloud_transformed = tf.matmul(point_cloud, transform)
19         input_image = tf.expand_dims(point_cloud_transformed, -1)
```

```
20     net = tf_util.conv2d(input_image, 64, [1,3],
21                          padding='VALID', stride=[1,1],
22                          bn=True, is_training=is_training,
23                          scope='conv1', bn_decay=bn_decay)
24     net = tf_util.conv2d(net, 64, [1,1],
25                          padding='VALID', stride=[1,1],
26                          bn=True, is_training=is_training,
27                          scope='conv2', bn_decay=bn_decay)
```

```
28     with tf.variable_scope('transform_net2') as sc:
29         transform = feature_transform_net(net, is_training, bn_decay, K=64)
30         end_points['transform'] = transform
31         net_transformed = tf.matmul(tf.squeeze(net, axis=[2]), transform)
32         net_transformed = tf.expand_dims(net_transformed, [2])
```

```
33     net = tf_util.conv2d(net_transformed, 64, [1,1],
34                          padding='VALID', stride=[1,1],
35                          bn=True, is_training=is_training,
36                          scope='conv3', bn_decay=bn_decay)
37     net = tf_util.conv2d(net, 128, [1,1],
38                          padding='VALID', stride=[1,1],
39                          bn=True, is_training=is_training,
40                          scope='conv4', bn_decay=bn_decay)
41     net = tf_util.conv2d(net, 1024, [1,1],
42                          padding='VALID', stride=[1,1],
43                          bn=True, is_training=is_training,
44                          scope='conv5', bn_decay=bn_decay)
```

```
55     # Symmetric function: max pooling
56     net = tf_util.max_pool2d(net, [num_point,1],
57                              padding='VALID', scope='maxpool')
58     net = tf.reshape(net, [batch_size, -1])
59     net = tf_util.fully_connected(net, 512, bn=True, is_training=is_training,
60                                   scope='fc1', bn_decay=bn_decay)
61     net = tf_util.dropout(net, keep_prob=0.7, is_training=is_training,
62                           scope='dp1')
63     net = tf_util.fully_connected(net, 256, bn=True, is_training=is_training,
64                                   scope='fc2', bn_decay=bn_decay)
65     net = tf_util.dropout(net, keep_prob=0.7, is_training=is_training,
66                           scope='dp2')
67     net = tf_util.fully_connected(net, 40, activation_fn=None, scope='fc3')
68     return net, end_points
```



Training code (train_one_epoch)

```
159 def train_one_epoch(sess, ops, train_writer):
160     """ ops: dict mapping from string to tf ops """
161     is_training = True
162
163     # Shuffle train files
164     train_file_idxs = np.arange(0, len(TRAIN_FILES))
165     np.random.shuffle(train_file_idxs)
166
167     for fn in range(len(TRAIN_FILES)):
168         log_string('-----' + str(fn) + '-----')
169         current_data, current_label = provider.loadDataFile(TRAIN_FILES[train_file_idxs[fn]])
170         current_data = current_data[:,0:NUM_POINT,:]
171         current_data, current_label, _ = provider.shuffle_data(current_data, np.squeeze(current_label))
172         current_label = np.squeeze(current_label)
173
174         file_size = current_data.shape[0]
175         num_batches = file_size // BATCH_SIZE
176
177         total_correct = 0
178         total_seen = 0
179         loss_sum = 0
180
181         for batch_idx in range(num_batches):
182             start_idx = batch_idx * BATCH_SIZE
183             end_idx = (batch_idx+1) * BATCH_SIZE
184
185             # Augment batched point clouds by rotation and jittering
186             rotated_data = provider.rotate_point_cloud(current_data[start_idx:end_idx, :, :])
187             jittered_data = provider.jitter_point_cloud(rotated_data)
188             feed_dict = {ops['pointclouds_pl']: jittered_data,
189                         ops['labels_pl']: current_label[start_idx:end_idx],
190                         ops['is_training_pl']: is_training,}
191             summary, step, _, loss_val, pred_val = sess.run([ops['merged'], ops['step'],
192                 ops['train_op'], ops['loss'], ops['pred']], feed_dict=feed_dict)
193             train_writer.add_summary(summary, step)
194             pred_val = np.argmax(pred_val, 1)
195             correct = np.sum(pred_val == current_label[start_idx:end_idx])
196             total_correct += correct
197             total_seen += BATCH_SIZE
198             loss_sum += loss_val
199
200         log_string('mean loss: %f' % (loss_sum / float(num_batches)))
201         log_string('accuracy: %f' % (total_correct / float(total_seen)))
```

Allow training of weights

Load training data

Train iteratively for training data

Data augmentation

Feed data and run session

Evaluation code (eval_one_epoch)

```
204 def eval_one_epoch(sess, ops, test_writer):
205     """ ops: dict mapping from string to tf ops """
206     is_training = False
207     total_correct = 0
208     total_seen = 0
209     loss_sum = 0
210     total_seen_class = [0 for _ in range(NUM_CLASSES)]
211     total_correct_class = [0 for _ in range(NUM_CLASSES)]
212
213     for fn in range(len(TEST_FILES)):
214         log_string('-----' + str(fn) + '-----')
215         current_data, current_label = provider.loadDataFile(TEST_FILES[fn])
216         current_data[:,0:NUM_POINT,:]
217         current_label = np.squeeze(current_label)
218
219         file_size = current_data.shape[0]
220         num_batches = file_size // BATCH_SIZE
221
222         for batch_idx in range(num_batches):
223             start_idx = batch_idx * BATCH_SIZE
224             end_idx = (batch_idx+1) * BATCH_SIZE
225
226             feed_dict = {ops['pointclouds_pl']: current_data[start_idx:end_idx, :, :],
227                         ops['labels_pl']: current_label[start_idx:end_idx],
228                         ops['is_training_pl']: is_training}
229             summary, step, loss_val, pred_val = sess.run([ops['merged'], ops['step'],
230                 ops['loss'], ops['pred']], feed_dict=feed_dict)
231             pred_val = np.argmax(pred_val, 1)
232             correct = np.sum(pred_val == current_label[start_idx:end_idx])
233             total_correct += correct
234             total_seen += BATCH_SIZE
235             loss_sum += (loss_val*BATCH_SIZE)
236             for i in range(start_idx, end_idx):
237                 l = current_label[i]
238                 total_seen_class[l] += 1
239                 total_correct_class[l] += (pred_val[i-start_idx] == l)
240
241     log_string('eval mean loss: %f' % (loss_sum / float(total_seen)))
242     log_string('eval accuracy: %f' % (total_correct / float(total_seen)))
243     log_string('eval avg class acc: %f' % (np.mean(np.array(total_correct_class)/np.array(total_seen_class,dtype=np.float))))
244
```

Freeze model weights

Loop the testing files

Load testing data


Feed data and run session

Some useful links

- ❑ MNIST example: https://github.com/aymericdamien/TensorFlow-Examples/blob/master/examples/3_NeuralNetworks/convolutional_network_raw.py
- ❑ PointNet example: <https://github.com/charlesq34/pointnet>
- ❑ More examples: <https://github.com/aymericdamien/TensorFlow-Examples>
- ❑ Point clouds projects: <https://github.com/Yochengliu/awesome-point-cloud-analysis>
- ❑ TensorFlow API: https://www.tensorflow.org/versions/r1.15/api_docs/python/tf
- ❑ Linux: <http://www.ee.surrey.ac.uk/Teaching/Unix/>
- ❑ Anaconda: <https://docs.conda.io/projects/conda/en/latest/user-guide/tasks/manage-environments.html>
- ❑ Stackoverflow: <https://stackoverflow.com/questions/tagged/tensorflow>

3D Object Detection

- Definition: It extends the concept of 2D object detection to the 3D physical space, and predicts the location, size, and orientation in the 3D space.




- Detection the target object in the image.
 - Predicting the location and size of each object in the image.
- Detection the target object in the 3D physical space.
 - Predicting the location, size, and orientation of each object.

3D Object Detection

- ❑ The 3D object detection is an important technique that can be widely applied in many realistic scenes for 3D perception.



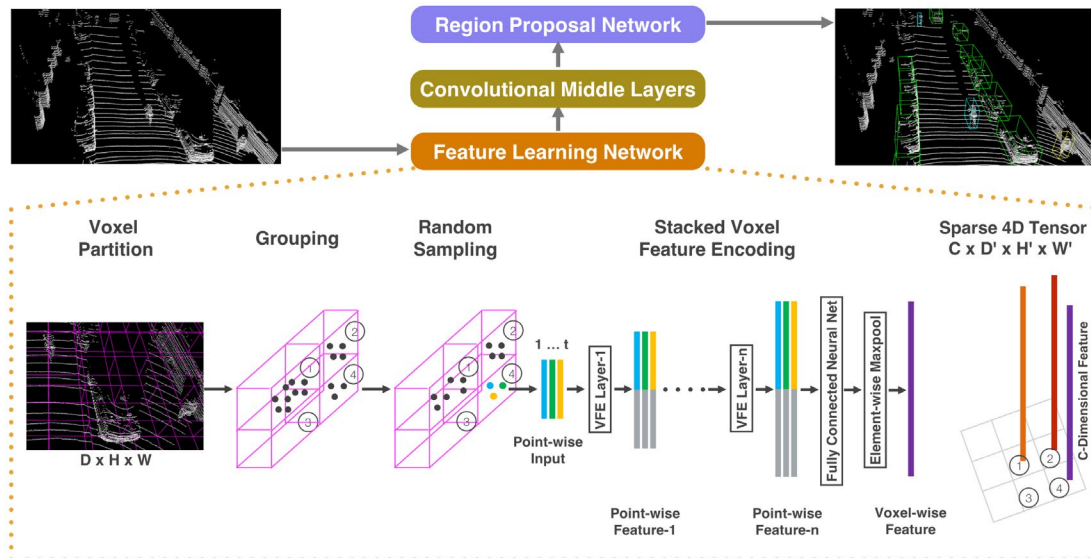
Autonomous Car



Indoor robots


VoxelNet (2017)

- The first work to apply PointNet to voxels, the overall process:
 - 1. Divide the space into voxels, use the VFE network in the voxel to extract local point cloud features, and obtain a three-dimensional feature volume.
 - 2. Send the features to the 3D convolutional network to further improve the expression ability, and compress the final output in the vertical direction to obtain a 2D feature map.
 - 3. Send the 2D feature map to the RPN network to generate a 3D box prediction.



VFE network in VoxelNet


- The VFE (Voxel-Feature-Encoding) structure is used to extract point cloud features within voxels:
 - The structure is similar to **PointNet**: a fully connected layer (including BN and ReLU) is used to transform the original coordinates of each point, and then the global feature is obtained by maximum pooling. Finally, the global feature is spliced onto the point-by-point feature as output.
 - **Stackable**: The output of VFE can be used as the input of the next VFE module, and multi-layer stacking improves the feature expression capability



The original input contains the original 3D coordinates of the point, the local 3D coordinates relative to the voxel center, and the reflectivity, a total of 7 dimensions.


Middle layers in VoxelNet

- 1. The feature volume output by the VFE layer has a size of $128 \times 10 \times 400 \times 352$ (feature dimension \times vertical \times horizontal \times depth)
- 2. It is fed into several 3D convolutional networks, and the vertical direction is downsampled by step size, and the output feature size is $64 \times 2 \times 400 \times 352$
- 3. The features of the two voxels in the vertical direction are concatenated to obtain a 2D feature map with a size of $128 \times 400 \times 352$ (feature dimension \times horizontal \times depth)



Region proposal network architecture

- ❑ VoxelNet uses a custom multi-scale detection network to give the final 3D bounding box prediction based on the 2D feature map
 - 1: The features output by the intermediate layer are fed into a U-Net-like structure to generate feature maps
 - 2. Based on the detection head of the wrong frame, two chain frames are set, so there are 2 category predictions and 14 bounding box regressors



Loss function

- The classification uses the two-class cross entropy BCE loss function, and the regression uses the SmoothL1 loss function.
 - For the regression loss, it follows the 2D detection R-CNN method to encode the residual between the true value box/prediction box and the anchor box.
 - The angle residual = the difference between the true value box and the anchor box angle.

$$L = \underbrace{\alpha \frac{1}{N_{\text{pos}}} \sum_i L_{\text{cls}}(p_i^{\text{pos}}, 1)}_{\text{Classification loss of positive samples}} + \underbrace{\beta \frac{1}{N_{\text{neg}}} \sum_j L_{\text{cls}}(p_j^{\text{neg}}, 0)}_{\text{Classification loss of negative samples}} + \underbrace{\frac{1}{N_{\text{pos}}} \sum_i L_{\text{reg}}(\mathbf{u}_i, \mathbf{u}_i^*)}_{\text{Regression loss of positive samples}}$$

$$\begin{aligned} \Delta x &= \frac{x_c^g - x_c^a}{d^a}, \Delta y = \frac{y_c^g - y_c^a}{d^a}, \Delta z = \frac{z_c^g - z_c^a}{h^a}, \\ \Delta l &= \log\left(\frac{l^g}{l^a}\right), \Delta w = \log\left(\frac{w^g}{w^a}\right), \Delta h = \log\left(\frac{h^g}{h^a}\right), \\ \Delta \theta &= \theta^g - \theta^a \end{aligned}$$
