

CS6382 Algorithm Analysis and Game Theory

Game Theory in Scheduling

Classic Scheduling Problems

- Set of m parallel machines M
- Set of n jobs J : each job j has (in most general case)
 - Arrival time r_j
 - Processing time (size) p_{ij} on machine $i \in M$
 - Value v_j
 - Deadline d_j

Classic Scheduling Problems

- A job j is considered **completed** if there exists machine i such that job j is processed for p_{ij} time units on i during $[a_j, d_j]$.
 - Let C_j denote the completion time of job j .
 - $C_j = \infty$ if job j is not completed.
- Preemptive vs. Non-preemptive
 - Whether job j is processed **continuously** on machine i

Objective of the Scheduler

- Schedule all jobs:
 - **Minimize the makespan** $\max_{j \in J} C_j$
 - Minimize total completion time $\sum_{j \in J} C_j$
 - Minimize total flow time $\sum_j (C_j - a_j)$
- Schedule some jobs to **maximize the total value** $\sum_{j: C_j \neq \infty} v_j$

Objective of the Machines

- Minimize the workload assigned to it.

Objective of the Machines

- Minimize the workload assigned to it.
- If there are payments for working:
 - maximize the profit: payment – workload

Objective of the Machines

- Minimize the workload assigned to it.
- If there are payments for working:
 - maximize the profit: payment – workload
- Suppose the machines are controlled by **selfish agents**:
 - Only machine i (agent i) knows p_{ij} , for all jobs $j \in J$
 - Machines may lie about p_{ij} to maximize their own utilities.

Objective of the Jobs

- Minimize its own completion time.
- Maximize its profit: v_j if completed; 0 otherwise.

Objective of the Jobs

- Minimize its own completion time.
- Maximize its profit: v_j if completed; 0 otherwise.
- If there is charging for being processed:
 - maximize the profit: value - charging

Objective of the Jobs

- Minimize its own completion time.
- Maximize its profit: v_j if completed; 0 otherwise.
- If there is charging for being processed:
 - maximize the profit: value - charging
- Suppose the jobs are controlled by **selfish agents**:
 - Only job j (agent j) knows a_j, v_j, d_j, p_{ij}
 - Jobs may lie about its information to maximize its own utility.

Game theory

- Compute a **Nash Equilibrium**: no agent has incentive to unilaterally change strategy.
- Design a scheduling mechanism under which **selfish machines** are truthful: misreporting is not beneficial.
- Design a scheduling mechanism under which **selfish jobs** are truthful: misreporting is not beneficial.

Nash Equilibrium

Nash Equilibrium

- Price of Anarchy (POA) : ratio between worst possible Nash equilibrium and the social optimum
 - All information are public.
 - How much worse will it be, if agents behave selfishly?
- Congestion Model
- Sequencing Model
- ...

Congestion model

- Selfish job chooses machine
- Utility: minimize completion time on the chosen machine.
- Objective of scheduler: minimize the makespan

Classic Results on Congestion Model

- On identical machines
 - $\text{POA} = \frac{3}{2}$ for 2 machines [Koutsoupias & Papadimitriou, STACS 1999].
 - $\text{POA} = \Omega\left(\frac{\log m}{\log \log m}\right)$ for m machines [Koutsoupias & Papadimitriou, STACS 1999]
 - $\text{POA} = O\left(\frac{\log m}{\log \log m}\right)$ for m machines [Czumaj & Vocking, SODA 2002]
- On related machines:
 - $\text{POA} \geq \Phi \approx 1.618$ on 2 machines [Koutsoupias & Papadimitriou, STACS 1999].
 - $\text{POA} = \Theta\left(\frac{\log m}{\log \log \log m}\right)$ for m machines [Czumaj & Vocking, SODA 2002]

Sequencing Model

- Utility of jobs: their own completion time
- Objective: minimize the makespan
- Need to define how jobs on the same machine are scheduled

Classic Results on Sequencing Model

- Shortest Processing Time first (SPT)
 - $2 - \frac{2}{m+1}$ for identical machines; $\Theta(\log m)$ on related machines; $O(m)$ for unrelated machines. [Immorlica et al., WINE 2005]
- Longest Processing Time first (LPT)
 - $\frac{4}{3} - \frac{1}{3m}$ for identical machines. [Christodoulou et al., ICALP 2004]
 - $\frac{4}{3} - \frac{1}{3m} \leq \text{POA} \leq 2 - \frac{2}{m}$ for related machines. [Immorlica et al., WINE 2005]

Selfish Machines

Model

- There are n jobs to be completed ($a_j = 0$, no deadline or value).
- Every machine i claims a processing time t_{ij} for processing job j .
- Scheduler assigns jobs to machines, and determines a payment for every machine, with the objective of minimizing makespan.
 - The actual processing time of j on i is given by p_{ij} , which might be different from t_{ij} .

Model

- Let $x_{ij} \in \{0,1\}$ be indicator of job j assigned to machine i
- Let $f_i(x)$ be the total payment to machine i .
- Utility of machine i : $f_i(x) - \sum_{j:x_{ij}=1} p_{ij}$
- Makespan: $\max_i \left\{ \sum_{j:x_{ij}=1} p_{ij} \right\}$

Truthful Mechanism Design

- If there is no payment:
 - All machines claim $t_{ij} = \infty$ to get rid of working.

Truthful Mechanism Design

- If there is no payment:
 - All machines claim $t_{ij} = \infty$ to get rid of working.
- Is there any scheduling algorithm + payment function that results in a truthful mechanism?

Truthful Mechanism Design

- If there is no payment:
 - All machines claim $t_{ij} = \infty$ to get rid of working.
- Is there any scheduling algorithm + payment function that results in a truthful mechanism?

YES!

Truthful Mechanism Design

- If there is no payment:
 - All machines claim $t_{ij} = \infty$ to get rid of working.
- Is there any scheduling algorithm + payment function that results in a truthful mechanism?
- A mechanism is truthful **if and only if** the scheduling algorithm is **Monotone**. [Saks & Yu, EC 2005]

Monotonicity

- For any (t_{i1}, \dots, t_{in}) and $(t'_{i1}, \dots, t'_{in})$ of machine i , let x and x' be the corresponding assignments, we have

$$\sum_{j=1}^m (x_{ij} - x'_{ij})(t_{ij} - t'_{ij}) \leq 0$$

- If job j is assigned to machines i , and we decrease only t_{ij} , then j should still be assigned to i .
 - Otherwise machine i has incentive to lie.

A Truthful Mechanism [Nisan & Ronen, STOC 1999]

- No truthful mechanism has approximation ratio less than 2, even on two machines.
 - Improved to $1 + \sqrt{2} \approx 2.414$ for $m \geq 3$ [Christodoulou et al., SODA 2007]
 - Improved to $1 + \Phi \approx 2.618$ for $m \rightarrow \infty$ [Koutsoupias & Vidali, MFCS 2007]
- There exists a truthful mechanism with approximation ratio m .
 - Optimal for two machines.

Lower Bound of 2

- Monotone \Leftrightarrow Truthful $\Rightarrow t_{ij} = p_{ij}$
- Fix any truthful mechanism on 2 machines, with 3 jobs:

Lower Bound of 2

- Monotone \Leftrightarrow Truthful $\Rightarrow t_{ij} = p_{ij}$
- Fix any truthful mechanism on 2 machines, with 3 jobs:
- If $p_{ij} = 1$ for all $i = 1, 2$ and $j = 1, 2, 3$, the scheduler (w.l.o.g.)
 - Schedule all jobs $\{1, 2, 3\}$ on the first machine, or
 - Schedule $\{1, 2\}$ on machine 1 and $\{3\}$ on machine 2.

Case 1: all jobs on machine **1**

1	$p_{11} = 1$	$p_{12} = 1$	$p_{13} = 1$
2			

Case 1: all jobs on machine 1

- By monotonicity:
- if machine 1 claims $t_{13} = 0$, without changing other t_{ij}
- then all jobs should still be assigned to machine 1

1	$p_{11} = 1$	$p_{12} = 1$	$p_{13} = 1$
2			

1	$p_{11} = 1$	$p_{12} = 1$	
2			$p_{13} = 0$

Case 1: all jobs on machine 1

- By monotonicity:
- if machine 1 claims $t_{13} = 0$, without changing other t_{ij}
- then all jobs should still be assigned to machine 1
- For the instance with $p_{13} = 0$, $p_{ij} = 1$ for other i, j :
- $\text{ALG} = 2$, $\text{OPT} = 1$

1	$p_{11} = 1$	$p_{12} = 1$	$p_{13} = 1$
2			

1	$p_{11} = 1$	$p_{12} = 1$	
2			$p_{13} = 0$

Case 2: jobs **{1,2}** on machine **1**

- If machine 2 claims $t_{23} = \epsilon, t_{21} = t_{22} = 1 + \epsilon$,
- by monotonicity: $(x'_{23} - 1) \cdot (1 - \epsilon) \geq \epsilon \cdot (x'_{21} + x'_{22})$

1	<table><tr><td>$p_{11} = 1$</td><td>$p_{12} = 1$</td></tr></table>	$p_{11} = 1$	$p_{12} = 1$
$p_{11} = 1$	$p_{12} = 1$		
2	<table><tr><td>$p_{23} = 1$</td></tr></table>	$p_{23} = 1$	
$p_{23} = 1$			

Case 2: jobs **{1,2}** on machine **1**

- If machine 2 claims $t_{23} = \epsilon, t_{21} = t_{22} = 1 + \epsilon$,
- by monotonicity: $0 \geq (x'_{23} - 1) \cdot (1 - \epsilon) \geq \epsilon \cdot (x'_{21} + x'_{22})$

1	<table><tr><td>$p_{11} = 1$</td><td>$p_{12} = 1$</td></tr></table>	$p_{11} = 1$	$p_{12} = 1$
$p_{11} = 1$	$p_{12} = 1$		
2	<table><tr><td>$p_{23} = 1$</td></tr></table>	$p_{23} = 1$	
$p_{23} = 1$			



Case 2: jobs **{1,2}** on machine **1**

- If machine 2 claims $t_{23} = \epsilon, t_{21} = t_{22} = 1 + \epsilon$,
- by monotonicity: $0 \geq (x'_{23} - 1) \cdot (1 - \epsilon) \geq \epsilon \cdot (x'_{21} + x'_{22}) \geq 0$.

1	<table><tr><td>$p_{11} = 1$</td><td>$p_{12} = 1$</td></tr></table>	$p_{11} = 1$	$p_{12} = 1$
$p_{11} = 1$	$p_{12} = 1$		
2	<table><tr><td>$p_{23} = 1$</td></tr></table>	$p_{23} = 1$	
$p_{23} = 1$			



Case 2: jobs **{1,2}** on machine **1**

- If machine 2 claims $t_{23} = \epsilon, t_{21} = t_{22} = 1 + \epsilon$,
- by monotonicity: $0 \geq (x'_{23} - 1) \cdot (1 - \epsilon) \geq \epsilon \cdot (x'_{21} + x'_{22}) \geq 0$.
- $\Rightarrow x'_{23} = 1, x'_{21} = x'_{22} = 0$: same schedule.



Case 2: jobs **{1,2}** on machine **1**

- If machine 2 claims $t_{23} = \epsilon, t_{21} = t_{22} = 1 + \epsilon$,
- by monotonicity: $0 \geq (x'_{23} - 1) \cdot (1 - \epsilon) \geq \epsilon \cdot (x'_{21} + x'_{22}) \geq 0$.
- $\Rightarrow x'_{23} = 1, x'_{21} = x'_{22} = 0$: same schedule.
- For the instance with $p_{23} = \epsilon, p_{21} = p_{22} = 1 + \epsilon$ for $\epsilon \rightarrow 0$:
- $\text{ALG} = 2, \text{OPT} = 1 + \epsilon$



Upper Bound of m

- Greedy: assign each job to the machine that processes it the fastest: $x_{ij} = 1$ if $t_{ij} < t_{lj}$ for all $l \neq i$.
 - Break tie by machine index

Upper Bound of m

- Greedy: assign each job to the machine that processes it the fastest: $x_{ij} = 1$ if $t_{ij} < t_{lj}$ for all $l \neq i$.
 - Break tie by machine index
- Monotone: $x_{ij}(t_{ij})$ is non-increasing.
 - Assignments of jobs are independent

Upper Bound of m

- Greedy: assign each job to the machine that processes it the fastest: $x_{ij} = 1$ if $t_{ij} < t_{lj}$ for all $l \neq i$.
 - Break tie by machine index
- Monotone: $x_{ij}(t_{ij})$ is non-increasing.
 - Assignments of jobs are independent
- Approximation ratio: $\text{ALG} \leq \sum_j \min_i \{p_{ij}\}$, $\text{OPT} \geq \frac{1}{m} \cdot \sum_j \min_i \{p_{ij}\}$

The payment function

- For every job j assigned to machine i , pay i the smallest claimed processing time from **other** machines.
 - Second Price Auction
- Truthful: for machine i and job j :
 - Suppose the payment is f_{ij} .
 - If $p_{ij} \geq f_{ij}$, then lying (to get the job) is not beneficial.
 - If $p_{ij} < f_{ij}$, then claiming any $t_{ij} < f_{ij}$ gives the same utility.

Randomized Truthful Mechanism

- Truthful in expectation: telling truth maximizes the expected utility
- Universally Truthful: for every realization of the random bits, telling truth maximizes the utility.
 - Since the schedule is random, the approximation ratio can be < 2
- For two machines: (universally truthful)
 - 1.75-approximation [Nisan & Ronen, STOC 1999]
 - Improved to 1.67 [Lu & Yu, STACS 2008]
 - Improved to 1.59 [Lu & Yu, WINE 2008]
- Lower bound of $2 - \frac{1}{m}$ [Muallem & Schapira, SODA 2007]

Selfish Related Machines

Model

- There are n jobs with $a_j = 0$ and $p_j > 0$, no deadline or value.
- Every machine i claims a speed t_i (truth: s_i):
 - processing job j requires $\frac{p_j}{t_i}$ units of time.

Model

- There are n jobs with $a_j = 0$ and $p_j > 0$, no deadline or value.
- Every machine i claims a speed t_i (truth: s_i):
 - processing job j requires $\frac{p_j}{t_i}$ units of time.
- Scheduler assigns jobs to machines, and determines payment f_i for machine i , with the objective of minimizing makespan.

Model

- There are n jobs with $a_j = 0$ and $p_j > 0$, no deadline or value.
- Every machine i claims a speed t_i (truth: s_i):
 - processing job j requires $\frac{p_j}{t_i}$ units of time.
- Scheduler assigns jobs to machines, and determines payment f_i for machine i , with the objective of minimizing makespan.
- Utility of machine i : $f_i(x) - \frac{1}{s_i} \cdot \sum_{j:x_{ij}=1} p_j$

Truthful Mechanisms

- Recall for unrelated machines: approximation: m ; hardness: 2 (conjectured to be m) [Nisan & Ronen, STOC 1999]

Truthful Mechanisms

- Recall for unrelated machines: approximation: m ; hardness: 2 (conjectured to be m) [Nisan & Ronen, STOC 1999]
- 3-approximate truthful (in expectation) [Archer & Tardos, FOCS 2001]
- PTAS, truthful (in expectation) [Dhangwatnotai et al., FOCS 2008]
- **Deterministic** truthful PTAS [Christodoulou & Kovacs, SODA 2010]
- Deterministic truthful PTAS for maximizing the minimum load, minimizing l_p norm [Epstein et al., SODA 2013]

Selfish Jobs

General Model

- m honest machines.
- Selfish jobs arrive **online**. For each job j :
 - Arrival time a_j
 - Deadline (departure time) d_j
 - Processing time (size) p_{ij} on machine i
 - Value v_j if it is completed
- Assumptions on selfish behaviors:
 - No earlier arrival, later deadline or shorter processing time.

Welfare Maximization Model


- m **identical** machines.
- Selfish jobs arrive **online**. For each job j :
 - Arrival time a_j , deadline d_j , size p_j , value v_j .
- Scheduler: maximize the **total value** of jobs completed.
- Selfish job j : maximize $q_j(x) \cdot v_j - f_j(x)$
 - $q_j(x) \in \{0,1\}$ indicates whether j is completed
 - $f_j(x)$: charging to job j for being processed
 - x : the schedule, which depends on the profiles of jobs

Unit-size Jobs

- Suppose $p_j = 1$ for all job j .
- Suppose a_j, d_j are integers.
- 2-competitive truthful mechanism [Hajiaghayi et al., EC 2005]
 - Greedy algorithm: at the beginning of every time slot, schedule the pending job with maximum value.
 - Greedy is monotone \Rightarrow truthful.
 - Greedy is 2-competitive.

Dominance and Monotonicity

- (a, d, v) dominates (a', d', v') if $a \leq a' \leq d' \leq d$, and $v > v'$
 - Longer active time window, larger value



Dominance and Monotonicity

- (a, d, v) dominates (a', d', v') if $a \leq a' \leq d' \leq d$, and $v > v'$
 - Longer active time window, larger value
- **Monotone** scheduling algorithm: if job j is scheduled with profile (a', d', v') , then j will be scheduled with profile (a, d, v) .

Dominance and Monotonicity

- (a, d, v) dominates (a', d', v') if $a \leq a' \leq d' \leq d$, and $v > v'$
 - Longer active time window, larger value
- **Monotone** scheduling algorithm: if job j is scheduled with profile (a', d', v') , then j will be scheduled with profile (a, d, v) .
- Monotone \Rightarrow truthful (with a carefully defined payment function)

A Simple Payment Function

- For claimed profile (a, d, v) , payment function

$$p_j(a, d, v) = q_j(a, d, v) \cdot \min\{v' : q_j(a, d, v') = 1\}$$

- Payment = smallest value that gets the job done.

A Simple Payment Function

- For claimed profile (a, d, v) , payment function

$$p_j(a, d, v) = q_j(a, d, v) \cdot \min\{v' : q_j(a, d, v') = 1\}$$

- Payment = smallest value that gets the job done.
- For any a and d , telling truth on v is optimal.
- By monotonicity, lying about a and d is not beneficial.

A Simple Payment Function

- For claimed profile (a, d, v) , payment function

$$p_j(a, d, v) = q_j(a, d, v) \cdot \min\{v' : q_j(a, d, v') = 1\}$$

- Payment = smallest value that gets the job done.
- For any a and d , telling truth on v is optimal.
- By monotonicity, lying about a and d is not beneficial.
- Does not apply to randomized/fractional assignments.

General Payment Function

- For claimed profile (a, d, v) , payment function

$$f_j(a, d, v) = q_j(a, d, v) \cdot v - \int_0^v q_j(a, d, x) dx$$

General Payment Function

- For claimed profile (a, d, v) , payment function

$$f_j(a, d, v) = q_j(a, d, v) \cdot v - \int_0^v q_j(a, d, x) dx$$

- Truthful if and only if (a_j, d_j, v_j) maximizes utility

$$q_j(a, d, v) \cdot v_j - f_j(a, d, v)$$

Suppose otherwise

- Suppose exist $a \geq a_j, d \leq d_j$ and v such that

$$q_j(a, d, v) \cdot v_j - f_j(a, d, v) > q_j(a_j, d_j, v_j) \cdot v_j - f_j(a_j, d_j, v_j)$$

Suppose otherwise

- Suppose exist $a \geq a_j, d \leq d_j$ and v such that

$$q_j(a, d, v) \cdot v_j - f_j(a, d, v) > q_j(a_j, d_j, v_j) \cdot v_j - f_j(a_j, d_j, v_j)$$

$$f_j(a_j, d_j, v_j) = q_j(a_j, d_j, v_j) \cdot v_j - \int_0^{v_j} q_j(a_j, d_j, x) dx$$

Suppose otherwise

- Suppose exist $a \geq a_j, d \leq d_j$ and v such that

$$q_j(a, d, v) \cdot v_j - f_j(a, d, v) > \int_0^{v_j} q_j(a_j, d_j, x) dx$$

Suppose otherwise

- Suppose exist $a \geq a_j, d \leq d_j$ and v such that

$$q_j(a, d, v) \cdot v_j - f_j(a, d, v) > \int_0^{v_j} q_j(a_j, d_j, x) dx$$

$$f_j(a, d, v) = q_j(a, d, v) \cdot v - \int_0^v q_j(a, d, x) dx$$

Suppose otherwise

- Suppose exist $a \geq a_j, d \leq d_j$ and v such that

$$q_j(a, d, v) \cdot (v_j - v) + \int_0^v q_j(a, d, x) dx > \int_0^{v_j} q_j(a_j, d_j, x) dx$$

Suppose otherwise

- Suppose exist $a \geq a_j, d \leq d_j$ and v such that

$$q_j(a, d, v) \cdot (v_j - v) + \int_0^v q_j(a, d, x) dx > \int_0^{v_j} q_j(a_j, d_j, x) dx$$

Consider three cases: $v = v_j$, $v < v_j$ and $v > v_j$.

If $v = v_j$

- Suppose exist $a \geq a_j, d \leq d_j, v = v_j$ such that

$$q_j(a, d, v) \cdot (v_j - v) + \int_0^v q_j(a, d, x) dx > \int_0^{v_j} q_j(a_j, d_j, x) dx$$

If $v = v_j$

- Suppose exist $a \geq a_j, d \leq d_j, v = v_j$ such that

$$q_j(a, d, v) \cdot (v_j - v) + \int_0^v q_j(a, d, x) dx > \int_0^{v_j} q_j(a_j, d_j, x) dx$$

$$\int_0^v q_j(a, d, x) dx > \int_0^{v_j} q_j(a_j, d_j, x) dx$$

If $v = v_j$

- Suppose exist $a \geq a_j, d \leq d_j, v = v_j$ such that

$$q_j(a, d, v) \cdot (v_j - v) + \int_0^v q_j(a, d, x) dx > \int_0^{v_j} q_j(a_j, d_j, x) dx$$

$$\int_0^v q_j(a, d, x) dx > \int_0^{v_j} q_j(a_j, d_j, x) dx$$

(a_j, d_j, x) dominates (a, d, x) : $q_j(a_j, d_j, x) \geq q_j(a, d, x)$
--

If $v < v_j$

- Suppose exist $a \geq a_j, d \leq d_j, v < v_j$ such that

$$q_j(a, d, v) \cdot (v_j - v) + \int_0^v q_j(a, d, x) dx > \int_0^{v_j} q_j(a_j, d_j, x) dx$$

If $v < v_j$

- Suppose exist $a \geq a_j, d \leq d_j, v < v_j$ such that

$$q_j(a, d, v) \cdot (v_j - v) + \int_0^v q_j(a, d, x) dx > \int_0^{v_j} q_j(a_j, d_j, x) dx$$

$$q_j(a, d, v) \cdot (v_j - v) > \int_v^{v_j} q_j(a_j, d_j, x) dx$$

If $v < v_j$

- Suppose exist $a \geq a_j, d \leq d_j, v < v_j$ such that

$$q_j(a, d, v) \cdot (v_j - v) + \int_0^v q_j(a, d, x) dx > \int_0^{v_j} q_j(a_j, d_j, x) dx$$

$$q_j(a, d, v) \cdot (v_j - v) > \int_v^{v_j} q_j(a, d, x) dx$$

If $v < v_j$

- Suppose exist $a \geq a_j, d \leq d_j, v < v_j$ such that

$$q_j(a, d, v) \cdot (v_j - v) + \int_0^v q_j(a, d, x) dx > \int_0^{v_j} q_j(a_j, d_j, x) dx$$

$$q_j(a, d, v) \cdot (v_j - v) > \int_v^{v_j} q_j(a, d, x) dx$$

(a, d, x) dominates (a, d, v) : $q_j(a, d, x) \geq q_j(a, d, v)$
--

If $v > v_j$

- Suppose exist $a \geq a_j, d \leq d_j, v > v_j$ such that

$$q_j(a, d, v) \cdot (v_j - v) + \int_0^v q_j(a, d, x) dx > \int_0^{v_j} q_j(a_j, d_j, x) dx$$

If $v > v_j$

- Suppose exist $a \geq a_j, d \leq d_j, v > v_j$ such that

$$q_j(a, d, v) \cdot (v_j - v) + \int_0^v q_j(a, d, x) dx > \int_0^{v_j} q_j(a_j, d_j, x) dx$$

$$\int_{v_j}^v q_j(a, d, x) dx > \int_0^v q_j(a, d, x) dx - \int_0^{v_j} q_j(a, d, x) dx > q_j(a, d, v) \cdot (v - v_j)$$

(a, d, v) dominates (a, d, x) : $q_j(a, d, v) \geq q_j(a, d, x)$
--

General Payment Function

- For claimed profile (a, d, v) , payment function

$$f_j(a, d, v) = q_j(a, d, v) \cdot v - \int_0^v q_j(a, d, x) dx$$

- **Truthful** since (a_j, d_j, v_j) maximizes utility

$$q_j(a, d, v) \cdot v_j - f_j(a, d, v)$$

Greedy algorithm is

- Monotone: if not completed, then remains not completed with a later arrival time, earlier deadline and smaller value.

Greedy algorithm is

- Monotone: if not completed, then remains not completed with a later arrival time, earlier deadline and smaller value.
- 2-Competitive: let A be the set of jobs completed by Greedy.
 - Consider any job $j \in \text{OPT} \setminus A$.
 - Let t_j be the time slot j is processed in OPT .
 - In Greedy, the job processed at t_j has higher value.
 - $v(\text{OPT}) \leq v(A) + v(\text{OPT} \setminus A) \leq 2 \cdot v(A)$

Unit-size Jobs (Asynchronous)

- Suppose $p_j = 1$ for all job j .
- Suppose a_j, d_j are non-negative reals.
- Restart/preemption is necessary:
 - Job with tight deadline and large value arrive when all machines are busy

Unit-size Jobs (Asynchronous)

- Suppose $p_j = 1$ for all job j .
- Suppose a_j, d_j are non-negative reals.
- Restart/preemption is necessary:
 - Job with tight deadline and large value arrive when all machines are busy
- 5-competitive truthful mechanism with restart [Hajiaghayi et al., EC 2005]

Unit-size Jobs (Asynchronous)

- Suppose $p_j = 1$ for all job j .
- Suppose a_j, d_j are non-negative reals.
- Restart/preemption is necessary:
 - Job with tight deadline and large value arrive when all machines are busy
- 5-competitive truthful mechanism with restart [Hajiaghayi et al., EC 2005]
- Lower bound of 5 for any restart/preemptive truthful mechanism [Chen et al. IJCAI 2017]

Arbitrary Size Jobs

- Job j has arrival time a_j , deadline d_j , size p_j , and value v_j
- $\omega(1)$ hardness for online preemptive algorithms, when jobs have tight deadlines: $d_j = a_j + p_j$. [Canetti & Irani, SICOMP 1998]
- $O(1)$ -competitive preemptive algorithm if the input instance has slack $s = \min_j \left\{ \frac{d_j - a_j}{p_j} \right\} > 1$. [Lucier et al., SPAA 2013]

Arbitrary Size Jobs

- Truthful $O(1)$ -competitive algorithm [Azar et al., EC 2015]
 - Slack $s = \min_j \left\{ \frac{d_j - a_j}{p_j} \right\} > 1$.
 - Monotone \Rightarrow truthful
 - β -Committed (by time $d_j - \beta \cdot p_j$, inform a job j whether it can be completed), given that $s \geq s(\beta)$ is sufficiently large.

Makespan Minimization Model

- Every selfish job j :
 - no deadline or value
 - minimize its own completion time
- Scheduler: minimize the makespan
 - Identical machines
 - Related machines
 - Unrelated machines

Truthful Mechanisms

- Any truthful mechanism without payments?

Truthful Mechanisms

- Any truthful mechanism without payments? **Greedy!**
 - Whenever a job arrives, schedule it to the machine that minimizes its completion time.
 - Break tie (of arrival time) by identity of jobs.
 - Truthful since job utilities are maximized.

Truthful Mechanisms

- Any truthful mechanism without payments? **Greedy!**
 - Whenever a job arrives, schedule it to the machine that minimizes its completion time.
 - Break tie (of arrival time) by identity of jobs.
 - Truthful since job utilities are maximized.
- Competitive ratio of Greedy
 - Identical machines: 2 [Graham, Bell System Technical Journal 1966]
 - Related machines: $\Theta(\log m)$ [Aspnes et al., STOC 1993]
 - Unrelated machines: $\Theta(m)$ [Aspnes et al., STOC 1993]

Better Competitive Ratios?

- Posted price mechanisms:
 - Before the next job arrive, set a price for every machine.
 - Each job decides which machine it goes to, and pays the price.
 - Truthful since jobs make their own decisions.
 - Static price: fixed price for each machine.
 - Dynamic price: prices change depending on current schedule.

Better Competitive Ratios?

- Posted price mechanisms:
 - Before the next job arrive, set a price for every machine
 - Each job decides which machine it goes to, and pays the price
 - Truthful since jobs make their own decisions.
 - Static price: fixed price for each machine
 - Dynamic price: prices change depending on current schedule
- On related machines: $O(1)$ -competitive [Feldman et al., EC 2017]
 - By dynamic prices

Constant Competitive Algorithm

- Online algorithm Slow-Fit [Azar et al., Journal of Algorithms 1997]
 - Maintain an upper bound Λ on **OPT**
 - Schedule every job to the slowest machine without violating the upper bound $2 \cdot \Lambda$.
 - Double Λ otherwise.
 - $O(1)$ -competitive.
 - Not truthful

Constant Competitive Algorithm

- Online algorithm Slow-Fit [Azar et al., Journal of Algorithms 1997]
 - Maintain an upper bound Λ on **OPT**
 - Schedule every job to the slowest machine without violating the upper bound $2 \cdot \Lambda$.
 - Double Λ otherwise.
 - $O(1)$ -competitive.
 - Not truthful
- How can we make Slow-Fit truthful?

Constant Competitive Algorithm

- Online algorithm Slow-Fit [Azar et al., Journal of Algorithms 1997]
 - Maintain an upper bound Λ on **OPT**
 - Schedule every job to the slowest machine without violating the upper bound $2 \cdot \Lambda$.
 - Double Λ otherwise.
 - $O(1)$ -competitive.
 - Not truthful
- How can we make Slow-Fit truthful? By dynamic prices.

An illustrating example [Feldman et al., EC 2017]

- Two machines with speeds $s_1 < s_2$
- Suppose currently, machine 1 finishes its jobs at time l_1 ; machine 2 finishes its jobs at time l_2 .
- Current upper bound: Λ

An illustrating example [Feldman et al., EC 2017]

- When a job with processing time p arrives
 - Slow-Fit assigns the job to machines 1 if $l_1 + \frac{p}{s_1} \leq 2 \cdot \Lambda$
 - Set price π (**independent of p**) on machine 2 such that

$$l_1 + \frac{p}{s_1} \leq l_2 + \frac{p}{s_2} + \pi \Leftrightarrow l_1 + \frac{p}{s_1} \leq 2 \cdot \Lambda$$

- Set $\pi = l_1 - l_2 + \left(1 - \frac{s_1}{s_2}\right) \cdot (2 \cdot \Lambda - l_1)$

General Algorithm

- By dynamic prices on machines:
- Selfish behavior of jobs results in Slow-Fit?

General Algorithm

- By dynamic prices on machines:
- Selfish behavior of jobs results in ~~Slow-Fit?~~



General Algorithm

- By dynamic prices on machines:
- Selfish behavior of jobs results in ~~Slow-Fit?~~
- Selfish behavior of jobs results in **Flex-Fit**, an algorithm similar to Slow-Fit, but has a higher flexibility on job assignments.
 - The competitive ratio of Flex-Fit is only a constant times larger than that of Slow-Fit. [Feldman et al., EC 2017]

Selfish Machines and Selfish Jobs

Classic Online Load Balancing Problem

- M : a set of m related machines, each machine i has speed s_i ;
- N : a set of n online jobs, each job j has size p_j ;
- Solution: Process all jobs on the machines without preemption and migration;
- Goal: Minimize the maximum finishing time over all machines (makespan).




Classic Online Load Balancing Problem

- M : a set of m related **machines**, each machine i has speed s_i ;
- N : a set of n online **jobs**, each job j has size p_j ;
- Solution: Process all jobs on the machines **without preemption and migration**;
- Goal: Minimize the maximum finishing time over all machines (makespan).
- Has a long history [Graham69, FW00, BCK00, ...]

Facts in Multi-agent System


1. The **jobs** are controlled by strategic agents:
 - The size is only known to the agent and her goal is to **minimize her own cost** (finishing time + potential payment).
 - Truthful mechanisms against strategic jobs [AAF'97, FFR'17, ...]



Facts in Multi-agent System


2. The **machines** are controlled by strategic agents:

- From the view of **incentive compatibility**:
The assignment and payment should be **truthful**.
[AAS'05, Kovacs'05, CK'13, ELS'16]
- From the view of **stability**:
The assignment should be **fair**.
[DPR'11, SR'13]



Our Problem

- Both jobs and machines are controlled by strategic agents.



Target Mechanisms

- **Truthful** against strategic jobs:
 - It is each job's best strategy to report her true size.
- **Well-behaved** for machines (\approx weighted Proportionality):
 - A machine has workload no smaller than that of any slower machine.
- **Why well-behaved mechanisms?**
 - Guarantees the fairness among the machines (stable cooperation);
 - Implies truthfulness of machines in some cases.
- **Benchmark:** Optimal offline makespan.

To what extent the makespan can be minimized under a well behaved mechanism that is truthful for strategic jobs?

Hardness of well-behavior

Hardness of well-behaved schedules

Theorem 1. *Well behaved schedules have competitive ratios $\Omega(\sqrt{m})$.*

- “Proof”:

- A set of m machines with speeds:

$$s_1 = 1 + \frac{1}{2^{m-1}},$$

$$s_2 = 1 + \frac{1}{2^{m-2}},$$


\vdots

$$s_m = 2.$$

- A set of m online jobs with sizes $p_j = s_j$.

- $OPT = 1$ by assigning job i to machine i .

- **Bad case:** the jobs arrive in increasing order of sizes.



Our Mechanism

Almost Well-behaved Algorithms

- Strong lower bound \rightarrow relax well-behavior
- γ -Almost well-behavior: a machine i has no smaller workload than machine i' only if $s_i \geq \gamma s_{i'}$.
- In our paper, we set $\gamma = 2$.
- Preprocessing:
 - Rounding each speed s_i down to the largest power of 2.
 - $s_1 \leq s_2 \leq \dots \leq s_n$ and each s_i is some power of 2.
- Next:
 - Design a well-behaved mechanism on rounded speeds that is truthful to jobs.



Dynamic Posted Pricing Mechanism

- Posted Prices of Machines:



Machine i : $price_i$ and current $workload_i$

The job can select the machine which **minimizes her cost** (finishing time + payment).



Dynamic Posted Pricing Mechanism

- Posted Prices of Machines:
 - **Inherently truthful**: if the prices are determined irrespective of the size of the next coming job;
 - **Simple**:
 - Jobs need not trust nor understand the logic underlying the mechanism
 - Jobs are not required to reveal their sizes (agents may not be aware of her true type).
- **Difficulty**:
 - Designing prices such that the selections of jobs imply well-behavior.

Dynamic Posted Pricing Mechanism, cont.

Multiple machines with the same speed?

Set the prices to be infinity except the one with minimum makespan.


- Assume all speeds s_i are different: $s_1 < s_2 < \dots < s_n$.
- In each round, a job j arrives:
 - Let $C_i(j)$ be the makespan of machine i before j arrives.
 - Let $\pi_m = 0$, and for each $i < m$, $\pi_i := \frac{s_i}{s_{i+1}} (C_{i+1}(j) - C_i(j))$.
 - Let $\rho_i(j) = \sum_{l \geq i} \pi_l$ be the price of machine i before j arrives.
 - So j will select the machine i such that $C_i(j) + \frac{p_i}{s_i} + \rho_i(j)$ is minimized.
 - Resulting schedule is well-behaved.

Theorem 2. *The mechanism is $O(\log m)$ -competitive in general and $O(1)$ -competitive if the sizes of online jobs are bounded.*

Corollaries

Corollary. *Combining Myerson's Lemma, our mechanism is also truthful for the machines if*

1. *all jobs have unit size; or*
2. *$m = 2$, i.e., two machines.*



Conclusion

Conclusion

- We show that any well-behaved mechanism must have a competitive ratio $\Omega(\sqrt{m})$.
- We propose an almost well-behaved **dynamic posted-price mechanism** that is truthful for the jobs and has a competitive ratio of $O(\log m)$.
- We show that when all jobs have **unit size** or there are only **two machines**, our mechanism is actually **truthful** against strategic machines.

Future Direction

Truthful mechanism for both the **strategic jobs** and **strategic machines**.

- A naïve mechanism (VCG):
 - Let $s_1 \leq \dots \leq s_{m-1} \leq s_m$ be the reported speeds;
 - Machine m gets all jobs with payment $\frac{1}{p_{m-1}} \sum_{j \in J} s_j$;
 - $ALG \leq m \cdot OPT$.
- We would like to see better mechanisms.